



```
In [18]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.preprocessing import StandardScaler
```

```
In [19]: data = pd.read_csv("diabetes[1].csv")
```

## DATA CLEANING AND UNDERSTANDING

```
In [20]: print("◆ Shape of the dataset:", data.shape)
print("\n◆ Columns:\n", data.columns)
print("\n◆ Data Types:\n", data.dtypes)
print("\n◆ First 5 rows:\n", data.head())
```

- ◆ Shape of the dataset: (768, 9)
- ◆ Columns:  
Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],  
 dtype='object')
- ◆ Data Types:  

|                          |         |
|--------------------------|---------|
| Pregnancies              | int64   |
| Glucose                  | int64   |
| BloodPressure            | int64   |
| SkinThickness            | int64   |
| Insulin                  | int64   |
| BMI                      | float64 |
| DiabetesPedigreeFunction | float64 |
| Age                      | int64   |
| Outcome                  | int64   |

  
dtype: object
- ◆ First 5 rows:  

|   | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin | BMI  | \ |
|---|-------------|---------|---------------|---------------|---------|------|---|
| 0 | 6           | 148     | 72            | 35            | 0       | 33.6 |   |
| 1 | 1           | 85      | 66            | 29            | 0       | 26.6 |   |
| 2 | 8           | 183     | 64            | 0             | 0       | 23.3 |   |
| 3 | 1           | 89      | 66            | 23            | 94      | 28.1 |   |
| 4 | 0           | 137     | 40            | 35            | 168     | 43.1 |   |

|   | DiabetesPedigreeFunction | Age | Outcome |
|---|--------------------------|-----|---------|
| 0 | 0.627                    | 50  | 1       |
| 1 | 0.351                    | 31  | 0       |
| 2 | 0.672                    | 32  | 1       |
| 3 | 0.167                    | 21  | 0       |
| 4 | 2.288                    | 33  | 1       |

```
In [21]: print("\n◆ Missing Values per Column:\n", data.isnull().sum())
```

- ♦ Missing Values per Column:

|                          |   |
|--------------------------|---|
| Pregnancies              | 0 |
| Glucose                  | 0 |
| BloodPressure            | 0 |
| SkinThickness            | 0 |
| Insulin                  | 0 |
| BMI                      | 0 |
| DiabetesPedigreeFunction | 0 |
| Age                      | 0 |
| Outcome                  | 0 |

dtype: int64

```
In [22]: print("\n♦ Descriptive Statistics:", data.describe())
```

- ♦ Descriptive Statistics:

|       | Pregnancies | Glucose    | BloodPressure | SkinThickness | Insulin    | \ |
|-------|-------------|------------|---------------|---------------|------------|---|
| count | 768.000000  | 768.000000 | 768.000000    | 768.000000    | 768.000000 |   |
| mean  | 3.845052    | 120.894531 | 69.105469     | 20.536458     | 79.799479  |   |
| std   | 3.369578    | 31.972618  | 19.355807     | 15.952218     | 115.244002 |   |
| min   | 0.000000    | 0.000000   | 0.000000      | 0.000000      | 0.000000   |   |
| 25%   | 1.000000    | 99.000000  | 62.000000     | 0.000000      | 0.000000   |   |
| 50%   | 3.000000    | 117.000000 | 72.000000     | 23.000000     | 30.500000  |   |
| 75%   | 6.000000    | 140.250000 | 80.000000     | 32.000000     | 127.250000 |   |
| max   | 17.000000   | 199.000000 | 122.000000    | 99.000000     | 846.000000 |   |


  

|       | BMI        | DiabetesPedigreeFunction | Age        | Outcome    |
|-------|------------|--------------------------|------------|------------|
| count | 768.000000 | 768.000000               | 768.000000 | 768.000000 |
| mean  | 31.992578  | 0.471876                 | 33.240885  | 0.348958   |
| std   | 7.884160   | 0.331329                 | 11.760232  | 0.476951   |
| min   | 0.000000   | 0.078000                 | 21.000000  | 0.000000   |
| 25%   | 27.300000  | 0.243750                 | 24.000000  | 0.000000   |
| 50%   | 32.000000  | 0.372500                 | 29.000000  | 0.000000   |
| 75%   | 36.600000  | 0.626250                 | 41.000000  | 1.000000   |
| max   | 67.100000  | 2.420000                 | 81.000000  | 1.000000   |

```
In [23]: print("\n♦ Number of Duplicate Rows:", data.duplicated().sum())
```

- ♦ Number of Duplicate Rows: 0

```
In [50]: plt.figure(figsize=(15,8))
sns.boxplot(data=data)
plt.title("Boxplot for Outlier Detection")
plt.show()
```



### Simple Linear Regression

```
In [36]: data = pd.read_csv("diabetes[1].csv")
data['Age'] = data['Age'].replace(0, data['Age'].mean())
```

```
In [37]: X = data[['Age']]
y = data['Glucose']
```

```
In [38]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
In [39]: model = LinearRegression()
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
```

```
In [40]: mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print("Mean Squared Error:", mse)
print("R2 Score:", r2)
```

Mean Squared Error: 953.4242781436858


R<sup>2</sup> Score: 0.0525531542583737

```
In [59]: import matplotlib.pyplot as plt
import seaborn as sns

# Calculate the correlation matrix
correlation_matrix = data.corr()

# Set the figure size to make it small
plt.figure(figsize=(8, 6))
```

```
# Create the heatmap
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt='.2f', linewidths=1)
plt.title('Correlation Matrix of Features')
plt.show()
```



```
In [41]: plt.scatter(X_test, y_test, color='skyblue', label='Actual')
plt.plot(X_test, y_pred, color='red', linewidth=2, label='Regression Line')
plt.xlabel("Age")
plt.ylabel("Glucose")
plt.title("Simple Linear Regression: Age vs Glucose")
plt.legend()
plt.show()
```



```
In [56]: selected_columns_3 = ['DiabetesPedigreeFunction', 'Age', 'BMI']
sns.pairplot(data[selected_columns_3], height=1.5)
plt.show()
```

